

The Traveling Salesman Problem (TSP)

Overview

The Traveling Salesman Problem (TSP) is possibly *the* classic discrete optimization problem.

A preview :

- How is the TSP problem defined?
- What we know about the problem: NP-Completeness.
- The construction heuristics: Nearest-Neighbor, MST, Clarke-Wright, Christofides.
- K-OPT.
- Simulated annealing and Tabu search.
- The Held-Karp lower bound.
- Lin-Kernighan.
- Lin-Kernighan-Helsgaun.
- Exact methods using integer programming.



Our presentation will pull together material from various sources - see the references below. But most of it will come from [\[App12006\]](#), [\[John1997\]](#), [\[CS153\]](#).

Defining the TSP


The TSP is fairly easy to describe:

- Input: a collection of points (representing cities).
- Goal: find a tour of minimal length.
Length of tour = sum of inter-point distances along tour

Input:




A non-optimal tour:
A B E D C



The optimal tour:
A B C D E

- Details:
 - Input will be a list of n points, e.g., $(x_0, y_0), (x_1, y_1), \dots, (x_{n-1}, y_{n-1})$.
 - Solution space: all possible tours.
 - "Cost" of a tour: total length of tour.
→ sum of distances between points along tour
 - Goal: find the tour with minimal cost (length).

- Strictly speaking, we have defined the *Euclidean TSP*.
- There are really three kinds:
 - The Euclidean (points on the plane).
 - The *metric* TSP: triangle inequality is satisfied.
 - The *graph* TSP:



Goal: find a minimal-length tour among tours that only use edges in the graph

Exercise:

- For an n -point problem, what is the size of the solution space (i.e., how many possible tours are there)?
- What's an example of an instance that's metric but not Euclidean?

Some assumptions and notation for the remainder:

- Let $n = |V|$ = number of vertices.
- Euclidean version, unless otherwise stated.
→ Complete graph.

Some history

Early history:

- 1832: informal description of problem in German handbook for traveling salesmen.
- 1883 U.S. estimate: 200,000 traveling salesmen on the road
- 1850's onwards: circuit judges

Exercise: Find the following 14 cities in Illinois/Indiana on a map and identify the best tour you can:

Bloomington, Clinton, Danville, Decatur, Metamora, Monticello, Mt.Pulaski, Paris, Pekin, Shelbyville, Springfield, Sullivan, Taylorville, Urbana

- 1960's: Proctor and Gamble \$10K competition: a 33-city TSP.
→ Won by a CMU mathematician (and others).
- A related problem: the *Knight's tour*.
→ Start at bottom-left corner, and visit all squares exactly once and return to the start.

Exercise: Show how the Knight's tour can be converted into a TSP instance.

- The statisticians take an interest
→ What is the expected length of an optimal tour for uniformly-generated points in 2D?
 - Several early analytic estimates in the 1940's.
 - Famous Beardwood-Halton-Hammersley result [Bear1959]:
If L^* = optimal tour's length then $L^* / \sqrt{n} \rightarrow$ a constant β
 - β estimated to be 0.72 for unit-square.
- Human solutions:
 - To assess problem-solving skill.
 - Part of some neurological tests.

TSP's importance in computer science:

- TSP has played a starring role in the development of algorithms.
- Used as a test case for almost every new (discrete) optimization algorithm:
 - Branch-and-bound.
 - Integer and mixed-integer algorithms.
 - Local search algorithms.
 - Simulated annealing, Tabu, genetic algorithms.
 - DNA computing.

Some milestones:

- Best known optimal algorithm: Held-Karp algorithm in 1962, $O(n^2 2^n)$.
- Proof of NP-completeness: Richard Karp in 1972 [\[Karp1972\]](#).
→ Reduction from Vertex-Cover (which itself reduces from 3-SAT).
- Two directions for algorithm development:
 - Faster exact solution approaches (using linear programming).
→ Largest problem solved optimally: 85,900-city problem (in 2006).
 - Effective heuristics.
→ 1,904,711-city problem solved within 0.056% of optimal (in 2009)
- Optimal solutions take a long time
→ A 7397-city problem took three years of CPU time.
- Theoretical development: (let L_H = tour-length produced by heuristic, and let L^* be the optimal tour-length)
 - 1976: Sahni-Gonzalez result [\[Sahn1976\]](#). Unless $P=NP$ no polynomial-time TSP heuristic can guarantee $L_H/L^* \leq 2^{p(n)}$ for any fixed polynomial $p(n)$.
 - Various bounds on particular heuristics (see below).
 - 1992: Arora et al result [\[Aror1992\]](#). Unless $P=NP$, there exists $\epsilon > 0$ such that no polynomial-time TSP heuristic can guarantee $L_H/L^* \leq 1 + \epsilon$ for all instances satisfying the triangle inequality.
 - 1998: Arora result [\[Aror1998\]](#). For Euclidean TSP, there is an algorithm that is polyomial for fixed $\epsilon > 0$ such that $L_H^*/L^* \leq 1 + \epsilon$

Approximate solutions: nearest neighbor algorithm

Nearest-neighbor heuristic:


- Possibly the simplest to implement.
- Sometimes called Greedy in the literature.
- Algorithm:

```

1.  V = {1, ..., n-1}           // Vertices except for 0.
2.  U = {0}                     // Vertex 0.
3.  while V not empty
4.      u = most recently added vertex to U
5.      Find vertex v in V closest to u
6.      Add v to U and remove v from V.
7.  endwhile
8.  Output vertices in the order they were added to U

```

Exercise: What is the solution produced by Nearest-Neighbor for the following 4-point Euclidean TSP. Is it optimal?




- What we know about Nearest-Neighbor:
 - $L_H/L^* \leq O(\log n)$
 - There are instances for which $L_H/L^* = O(\log n)$
 - There are sub-classes of instances for which Nearest-Neighbor consistently produces the *worst* tour [\[Guti2007\]](#).


Approximate solutions: the Clarke-Wright heuristic

The Clarke-Wright algorithm: [\[Clar1964\]](#).


- The idea:
 - First identify a "hub" vertex:



- Compute starting cost as cost of going through hub:



- Identify "savings" for each pair of vertices:



- Take shortcuts and add them to final tour, as long as no cycles are created.

- Algorithm:


```

1.  Identify a hub vertex  $h$ 
2.   $V_H = V - \{h\}$ 
3.  for each  $i, j \neq h$ 
4.      compute  $\text{savings}(i, j)$ 
5.  endfor
6.  sortlist = Sort vertex pairs in decreasing order of savings
7.  while  $|V_H| > 2$ 
8.      try vertex pair  $(i, j)$  in sortlist order
9.      if  $(i, j)$  shortcut does not create a cycle
         and  $\text{degree}(v) \leq 2$  for all  $v$ 
10.         add  $(i, j)$  segment to partial tour
11.         if  $\text{degree}(i) = 2$ 
12.              $V_H = V_H - \{i\}$ 
13.         endif
14.         if  $\text{degree}(j) = 2$ 
15.              $V_H = V_H - \{j\}$ 
16.         endif
17.     endif
18. endwhile
19. Stitch together remaining two vertices and hub into final tour


```

- Example (from above):


- Suppose vertex 4 is the hub vertex:




- Suppose (2,3) provides the most savings:



- Next, (1,2) gets added
 - $\text{degree}(2) = 2$
 - must remove hub edge (2,4)




- Continuing ... let's say we obtain:



Only two vertices, 5 and 7, remain connected to the hub

- Finally, add last two vertices and hub into final tour:



Final tour


- What's known about the CW heuristic:
 - Bound is logarithmic: $L_H/L^* \leq O(\log n)$
 - Worst examples known: $L_H/L^* \geq O(\log(n) / \log\log(n))$

Approximate solutions: the MST heuristic


An approximation algorithm for (Euclidean) TSP that uses the MST: [\[Rose1977\]](#).

- The algorithm:
 - First find the minimum spanning tree (using any MST algorithm).
 - Pick any vertex to be the root of the tree.
 - Traverse the tree in *pre-order*.
 - Return the order of vertices visited in pre-order.


Exercise: What is the pre-order for this tree (starting at A)?



- Example:
 - Consider these 7 points:




◦ A minimum-spanning tree:




Minimum spanning tree

◦ Pick vertex A as the root:




◦ Traverse in pre-order:




Visit order: A E B D C F G


◦ Tour:



- Claim: the tour's length is no worse than twice the optimal tour's length.
 - Let L = the length of the tour produced by the algorithm.
 - Let L^* = the length of an optimal tour.
 - Let M = weight of the MST (total length).
 - Observe: if we remove any one edge from a tour, we will get a spanning tree.
 $\rightarrow L^* > M$.
 - Now consider a pre-order *tree walk* from the root, back to the root:



- Let W = length of this walk.
- Then, $W = 2M$ (each edge is traversed twice).
- Thus, $W < 2L^*$.
- Finally, we will show that $L \leq W$ and therefore, $L < 2L^*$.
- To see why, consider the tree walk from B to D:



$$\text{length BE} + \text{length ED} < \text{length BD}$$

$\rightarrow L$ takes a shorter route than W (triangle inequality).

- Thus, $L \leq W$.


What we know about this algorithm:

- The first heuristic to produce solutions within a constant of optimal.
- Easy to implement (since MST can be found efficiently).


Approximate solutions: the Christofides heuristic

The Christofides algorithm: [Chri1976].

- First, as background, we need to understand two things:
 - What is an Euler tour (for general graphs)?
 - A tour that traverses all edges exactly once (but may repeat vertices)




Euler tour exists




No Euler tour possible

- Famous result: a graph has an Euler tour if and only if all its vertices have even degree.
- What is a minimal matching for a given subset of vertices V' ?
 - A "best" (minimal weight) subset of edges with the property that no edges have a common vertex




V' = given set of vertices
for which matching is desired




a matching

- Important result: min-matching can be found in poly-time.
- The key ideas in the algorithm:
 - First find the MST




- Then identify the odd-degree vertices



- There are an *even* number of such odd-degree vertices.

Exercise: Why?


- Find a minimal matching of these odd-degree vertices and add those edges



Adding "match" edges to original graph may result in multiple edges between a vertex pair


Drawn as curved to distinguish from regular graph edges

- Now all vertices have even degree.
- Next, find an Euler tour.



An Euler tour that may revisit vertices


- Now, walk along in Euler tour, but skip visited nodes



- This produces a TSP tour.

An improved bound:

- We will show that $L_H/L^* \leq 1.5$
- Let M = cost of MST.
→ $L^* \geq M$ (as argued before).
- Note: we are performing a matching on an even number of vertices.
- Now consider the original odd-degree vertices



- Consider the optimal tour on just these (even # of) vertices.
- Let L_O = cost of this tour.
- Let e^1, e^2, \dots, e^{2k} be the edges.
- Note: $E_1 = \{e^1, e^3, \dots, e^{2k-1}\}$ is a matching.
- So is $E_2 = \{e^2, e^4, \dots, e^{2k}\}$
- Now at least one set has weight at most $L_O/2$.
→ Because both must add up to L_O .
- Also the optimal matching found earlier has less weight than either of these edge sets.
→ $\text{min-match-cost} \leq L_O/2 \leq L^*/2$.
- Thus $\text{min-match-cost} + M \leq L^* + L^*/2$
- But L_H uses edges (or shortcuts) from min-match and MST

$$\rightarrow L_H \leq L^* + L^*/2$$


Running time:

- Dominated by $O(n^3)$ time for matching.
- Best known matching algorithm: $O(n^{2.376})$

K-OPT

Constructive vs. local-search heuristics:

- All four heuristics above were *constructive*
→ A tour was built up step by step.
- In contrast, a local-search heuristic works as follows:




```

1. T = some starting tour      // Perhaps by using Christofides.
2. noChange = true
3. repeat
4.   T' = makeChangeTo (T)
5.   if T' < T
6.     T = T'
7.     noChange = false
8.   endif
9. until noChange
10. return T

```

2-OPT:

- Idea: Replace 2 edges and see if the cost improves.



- Find two edges and their endpoints.
- Swap endpoints.

- 2-OPT heuristic

```

1. T = some starting tour
2. noChange = true
3. repeat
4.   for all possible edge-pairs in T
5.     T' = tour by swapping end points in edge-pair
6.     if T' < T
7.       T = T'

```

```

8.         noChange = false
9.         break      // Quit loop as soon as an improvement is found
10.      endif
11.   endfor
12. until noChange
13. return T

```

- An alternative: find best tour with all possible swaps:

```

1. T = some starting tour
2. noChange = true
3. repeat
4.   Tbest = T
5.   for all possible edge-pairs in T
6.     T' = tour by swapping end points in edge-pair
7.     if T' < Tbest
8.       Tbest = T'
9.       noChange = false
10.    endif
11.  endfor
12.  T = Tbest
13. until noChange
14. return T

```

K-OPT:

- 3-OPT is what you can get by considering replacing 3 edges.
- K-OPT considers K edges.
- Each K-OPT can be time-consuming for $K > 3$.


What we know about K-OPT:

- For general graphs: $L_H/L^* \leq 0.25 n^{1/2k}$.
- For Euclidean case, $L_H/L^* \leq O(\log n)$.
- In practice: 2-OPT and 3-OPT are much better than the construction heuristics.
- Note: Any K-OPT move can be reduced to a sequence of 2-OPT moves.
→ But might it might require a long such sequence.


Local Optima and Problem Landscape

Local optima:

- Recall: greedy-local-search generates one state (tour) after another until no better neighbor can be found
→ does this mean the last one is optimal?
- Observe the trajectory of states:




- There is no guarantee that a greedy local search can find the (global) minimum.
- The last state found by greedy-local-search is a *local minimum*.
→ it is the "best" in its neighborhood.
- The *global minimum* is what we seek: the least-cost solution overall.
- The particular local minimum found by greedy-local-search depends on the start state:





Problem landscape:

- Consider TSP using a particular local-search algorithm:
 - Suppose we use a graph where the vertices represent states.
 - An edge is placed between two "neighbors"
e.g., for a 5-point TSP the neighbors of [0 1 2 3 4] are:



Neighbors of [0 1 2 3 4] using a 2-point swap

- The cost of each tour is represented as the "weight" of each vertex.
- Thus, a local-search algorithm "wanders" around this graph.
- Picture a 3D surface representing the cost *above* the graph.
→ this is the problem landscape for a particular problem and local-search algorithm.



- A large part of the difficulty in solving combinatorial optimization problems is the "weirdness" in landscapes
→ landscapes often have very little structure to exploit.
- Unlike continuous optimization problems, local shape in the landscape does NOT help point towards the global minimum.


Climbing out of local minima:

- A local-search algorithm gets "stuck" in a local minimum.
- One approach: re-run local-search many times with different starting points.
- Another approach (next): help a local-search algorithm "climb" out of local minima.


Tabu search

Key ideas: [\[Glov1990\]](#).

- Suppose we decide to climb out of local minima.
- Danger: could immediately return to same local minima.



- In tabu-search, you maintain a list of "tabu tours".
→ The algorithm avoids these.
- Each time you pick a minimum in a neighborhood, add that to the tabu list.




- Various alternatives to tabu-lists
 - Always add all neighborhood minimums.
 - Only add local minima.
- This way, Tabu forces more searching.
- A problem: a tabu-list can grow very long.
→ Need a *policy* for removing items, e.g.,
 - Least-recently used.
 - Throw out high-cost tours.


Simulated annealing


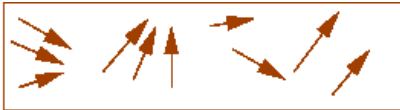


Background:

- What is *annealing*?
 - *Annealing* is a metallurgic process for improving the strength of metals.
 - Key idea: cool metal slowly during the forging process.
- Example: making bar magnets:
 - Wrong way to make a magnet:
 1. Heat metal bar to high temperature in magnetic field.




2. Cool rapidly (quench):



- Right way: cool slowly (anneal)
- Why slow-cooling works:
 - At high heat, magnetic dipoles are agitated and move around:
 
 - The magnetic field tries to force alignment:
 
 - If cooled rapidly, alignments tend to be less than optimal (local alignments):
 
 - With slow-cooling, alignments are closer to optimal (global alignment):
 

- Summary: slow-cooling helps because it gives molecules more time to "settle" into a globally optimal configuration.
- Relation between "energy" and "optimality"
 - The more aligned, the lower the system "energy".
 - If the dipoles are not aligned, some dipoles' fields will conflict with others.
 - If we (loosely) associate this "wasted" conflicting-fields with energy
 - better alignment is equivalent to lower energy.
 - Global minimum = lowest-energy state.
- The Boltzmann Distribution:
 - Consider a gas-molecule system (chamber with gas molecules):




- The state of the system is the particular snapshot (positions of molecules) at any time.
- There are high-energy states:



High-energy: molecules bunched up

and low-energy states:



Low energy state: molecules spread apart

- Suppose the states s_1, s_2, \dots have energies $E(s_1), E(s_2), \dots$
- A particular energy value E occurs with probability

$$P[E] = Z e^{-E/kT}$$

where Z and k are constants.

- Low-energy states are more probable at low temperatures:
 - Consider states s_1 and s_2 with energies $E(s_2) > E(s_1)$
 - The ratio of probabilities for these two states is:

$$r = P[E(s_1)]/P[E(s_2)] = e^{[E(s_2) - E(s_1)]/kT} = \exp([E(s_2) - E(s_1)]/kT)$$

Exercise : Consider the ratio of probabilities above:

- Question: what happens to r as T increases to infinity?
- Question: what happens to r as T decreases to zero?


What are the implications?

Key ideas in simulated annealing: [\[Kirk1983\]](#).

- Simulated annealing = a modified local-search.
- Use it to solve a combinatorial optimization problem.
- Associate "energy" with "cost".
 - Goal: find lowest-energy state.
- Recall problem with local-search: gets stuck at local minimum.
- Simulated annealing will allow jumps to higher-cost states.
- If randomly-selected neighbor has lower-cost, jump to it (like local-search does).
- If randomly-selected neighbor is of higher-cost
 - flip a coin to decide whether to jump to higher-cost state
 - Suppose current state is s with cost $C(s)$.
 - Suppose randomly-selected neighbor is s' with cost $C(s') > C(s)$.
 - Then, jump to it with probability

$$e^{-[C(s') - C(s)]/kT}$$

- Decrease coin-flip probability as time goes on:
 - by decreasing temperature T .
- Probability of jumping to higher-cost state depends on cost-difference:



Implementation:

- Pseudocode: (for TSP)

Algorithm: TSPSimulatedAnnealing (points)

Input: array of points

```

// Start with any tour, e.g., in input order
1. s = initial tour 0,1,...,n-1

// Record initial tour as best so far.
2. min = cost (s)
3. minTour = s

// Pick an initial temperature to allow "mobility"
4. T = selectInitialTemperature()

// Iterate "long enough"
5. for i=1 to large-enough-number
    // Randomly select a neighboring state.
6.    s' = randomNextState (s)
    // If it's better, then jump to it.
7.    if cost(s') < cost(s)
8.        s = s'
        // Record best so far:
9.        if cost(s') < min
10.           min = cost(s')
11.           minTour = s'
12.        endif
13.    else if expCoinFlip (s, s')
        // Jump to s' even if it's worse.
14.        s = s'
15.    endif // Else stay in current state.
    // Decrease temperature.
16.    T = newTemperature (T)
17. endfor

18. return minTour

```

Output: best tour found by algorithm

Algorithm: randomNextState (s)

Input: a tour s, an array of integers

```

// ... Swap a random pair of points ...

```

Output: a tour

Algorithm: expCoinFlip (s, s')

Input: two states s and s'

```

1.  p = exp ( -(cost(s') - cost(s)) / T)
2.  u = uniformRandom (0, 1)
3.  if u < p
4.      return true
5.  else
6.      return false

```

Output: true (if coinFlip resulted in heads) or false

- Implementation for other problems, e.g., BPP
 - The only thing that needs to change: define a `nextState` method for each new problem.
 - Also, some experimentation will be need for the temperature schedule.

Temperature issues:

- Initial temperature:
 - Need to pick an initial temperature that will accept large cost increases (initially).
 - One way:
 - Guess what the large cost increase might be.
 - Pick initial T to make the probability 0.95 (close to 1).
- Decreasing the temperature:
 - We need a *temperature schedule*.
 - Several standard approaches:
 - Multiplicative decrease: Use $T = a * T$, where a is a constant like 0.99.
 $\rightarrow T_n = a^n$.
 - Additive decrease: Use $T = T - a$, where a is a constant like 0.0001.
 - Inverse-log decrease: Use $T = a / \log(n)$.
 - In practice: need to experiment with different temperature schedules for a particular problem.

Analysis:

- How long do we run simulated annealing?
 - Typically, if the temperature becomes very, very small there's no point in further execution
 \rightarrow because probability of escaping a local minimum is miniscule.
- Unlike previous algorithms, there is no fixed running time.
- What can we say theoretically?
 - If the inverse-log schedule is used
 \rightarrow Can prove "probabilistic convergence to global minimum"
 \rightarrow Loosely, as the number of iterations increase, the probability of finding the global minimum tends to 1.

In practice:

- Advantages of simulated annealing:
 - Simple to implement.
 - Does not need much insight into problem structure.
 - Can produce reasonable solutions.
 - If greedy does well, so will annealing.
- Disadvantages:
 - Poor temperature schedule can prevent sufficient exploration of state space.
 - Can require some experimentation before getting it to work well.

- Precautions:
 - Always re-run with several (wildly) different starting solutions.
 - Always experiment with different temperature schedules.
 - Always pick an initial temperature to ensure high probability of accepting a high-cost jump.
 - If possible, try different neighborhood functions.
- Warning:
 - Just because it has an appealing origin, simulated annealing is not guaranteed to work
→ when it works, it's because it explores more of the state space than a greedy-local-search.
 - Simply running greedy-local-search on multiple starting points may be just as effective, and should be experimented with.

Variations:


- Use greedyNextState instead of the nextState function above.
 - Advantage: guaranteed to find local minima.
 - Disadvantage: may be difficult or impossible to climb out of a particular local minimum:
 - Suppose we are stuck at state s , a local minimum.
 - We probabilistically jump to s' , a higher-cost state.
 - When in s' , we will very likely jump back to s (unless a better state lies on the "other side").
 - Selecting a random next-state is more amenable to exploration.
→ but it may not find local minima easily.
- Hybrid nextState functions:
 - Instead of considering the entire neighborhood of 2-swaps, examine some fraction of the neighborhood.
 - Switch between different neighborhood functions during iteration.
- Maintain "tabu" lists:
 - To avoid jumping to states already seen before, maintain a list of "already-visited" states and exclude these from each neighborhood.
- Thermal cycling:
 - Periodically raise temperature and perform "re-starts".
 - The idea is to force more exploration of the state space.

The Held-Karp lower bound


Our presentation will follow the one in [\[Vale1997\]](#).

First, a definition:

- Consider a graph with vertices $\{1, \dots, n\}$:




- A 1-tree is a subgraph constructed as follows:




A 1-tree

=




a spanning tree
of vertices $\{2, \dots, n\}$

+




and the two cheapest
edges from vertex 1

- Temporarily remove vertex 1 (and its edges) and find a spanning tree for vertices $\{2, \dots, n\}$.
- Then pick add two cheapest edges from vertex 1.
- Note: every tour (including the optimal one) is a 1-tree.



- The *min-1-tree* is the lowest weighted 1-tree among all 1-trees.
→ This will be a lower bound for the optimal tour.
- A simple algorithm for the min-1-tree:
 - Find the MST for the graph without vertex 1.
 - Add the two cheapest edges from vertex 1.
- Is the min-1-tree a good bound?


Exercise: What is the difference between the optimal tour and the min-1-tree for this graph?



- The problem is: the MST can avoid using edges that the tour must take.

Held-Karp's idea:

- We will associate a π_i , a *vertex weight* with every vertex i .
- Define a modified graph G' as follows:
 - G' has the same vertices and edges as G .
 - Let e_{ij} = weight of edge (i,j) in G .
 - Let c_{ij} = weight of edge (i,j) in G' .
 - Then define $c_{ij} = e_{ij} + \pi_i + \pi_j$.
- For example:



Exercise: What is the difference between the min-1-tree and the optimal tour for the above modified graph G' ? What vertex weight for the top-right vertex best closes the gap between the min-1-tree and the optimal tour?

- Thus, one can *choose* weights so that the min-1-tree is as high as possible in G' .

In more detail:

- Let T be a 1-tree and T' be a tour.

- Let d_i^T = the degree of node i in T .
- Let $L(T, G)$ = cost of 1-tree T using graph G .
- Let $L(T', G)$ = cost of tour T' using graph G .
- Since every tour is a 1-tree, $\min_T L(T, G) \leq \min_{T'} L(T', G)$.
- Now, for a 1-tree T ,

$$L(T, G') = L(T, G) + \sum_{i \in T} (d_i^T - 2) \pi_i.$$
- Similarly, for a tour T' ,


$$L(T', G') = L(T', G) + \sum_{i \in T'} 2 \pi_i.$$
- Thus, subtracting and taking minimum, $\min_T L(T, G) + \sum_{i \in T} (d_i^T - 2) \pi_i \leq \min_{T'} L(T', G) = L^*$ (the optimal tour).
- To summarize, we want to find the min-1-tree with weights π and then correct for that by subtracting off the additional weights.
- Let $W(\pi) = \min_T L(T, G) + \sum_{i \in T} (d_i^T - 2) \pi_i$.
- Then, the desired "best" Held-Karp bound is: $\max_{\pi} W(\pi)$.

An optimization procedure:

- Let $V_{T(\pi)}$ be the vector (d_1^T, \dots, d_n^T) .
- Let $C_{T(\pi)}$ be the cost of min-1-tree using π .
- Then, write $W(\pi) = C_{T(\pi)} + \pi \cdot V_{T(\pi)}$.
- Next, suppose that π' is a vector in π -space such that $W(\pi') \geq W(\pi)$.
- Then, Held-Karp show that $(\pi' - \pi) \cdot V_{T(\pi)} \geq 0$.
- This means that larger values of $W(\pi')$ are in the right half-space pointed to by the vector $V_{T(\pi)}$.
- Next step: an iterative optimization procedure.


First, a little background on gradient-based optimization:

- Consider a (single-dimensional) function $f(x)$:



- Let $f'(x)$ denote the derivative of $f(x)$.
- The gradient at a point x is the value of $f'(x)$.
 \rightarrow Graphically, the slope of the tangent to the curve at x .

- Observe the following:



- To the left of the optimal value x^* , the gradient is negative.
- To the right, it's positive.
- We seek an iterative algorithm of the form

```

while not over
  if gradient < 0
    move rightwards
  else if gradient > 0
    move leftwards
  else
    stop // gradient = 0 (unlikely in practice, of course)
  endif
endwhile

```

- The gradient descent algorithm is exactly this idea:

```

while not over
  x = x - α f'(x)
endwhile

```

Here, we add a scaling factor α in case $f'(x)$ values are of a different order-of-magnitude:

Back to vertex-weight optimization:



- Unfortunately, we don't have a differentiable function.
- For this case, the Russian mathematician Polyak devised what's called the *sub-gradient* algorithm:
 - For a differentiable function, the gradient "points" in the right direction.
 - For a non-differentiable function, it's still possible to use a gradient that points in the right direction.
- For the vertex-weights, the iteration turns out to be: $\pi_i^{(m+1)} = \pi_i^{(m)} + \alpha^{(m)} (d_i - 2)$.
- Intuitively, this means:
 - Increase the weights for vertices with 1-min-tree degree > 2 .
 - Decrease the weights for vertices with 1-min-tree degree < 2 .
 - Thus, the iteration tries to force the 1-min-tree to be "tour-like".
- Polyak showed that sub-gradient iteration works if the stepsizes $\alpha^{(m)}$ are chosen properly:
 - $\alpha^{(m)} \rightarrow 0$
 - $\sum_m \alpha^{(m)} = \infty$
- To summarize:
 - Start with some vector of vertex-weights π .
 - Repeatedly apply the iteration $\pi_i^{(m+1)} = \pi_i^{(m)} + \text{stepsize} * \text{sub-gradient } V_{T(\pi)}$.

- Implementation issues:
 - Each iteration requires an MST computation.
 - Can be expensive for large n .
 - One approximation: reduce number of edges by considering only best k neighbors (e.g., $k=20$).

The Lin-Kernighan algorithm



Key ideas:

- Devised in 1973 by Shen Lin (co-author on BB(N) numbers) and Brian Kernighan (the "K" of K&R fame).
- Champion TSP heuristic 1973-89.
- LK is iterative:
 - Starts with a tour and repeatedly improves, until no improvement can be found.
- Idea 1: Make the K edges in K-OPT *contiguous*



- This is just the high-level idea
 - The algorithm actually alternates between a "current-tour-edge" and a "new-putative-edge".
- Let the K in K-OPT vary at each iteration.
 - Try to increase K gradually at each iteration.
 - Pick the best K (the best tour) along the way.
- Allow some limited backtracking.
- Use a tabu-list to create freshness in exploration.


Note: we will use an artificial depiction of a tour as follows:




This will be used to explain some ideas.

The LK algorithm in more detail:


- At each iteration, LK identifies a sequence of edges $x_1, y_1, x_2, y_2, \dots, x_k, y_k$ such that:



- Each x_i is an edge in the current tour.
 - Each y_i is NOT in the current tour.
 - They are all unique (no repetitions).
 - The last y_k returns to the starting point t_1
- We'll call this an *LK-move*.
 - For example:




- Notice that if we stop at any intermediate y_i , we get a 1-tree.




- Let G_1 = gain after first x-y-pair:

$$G_1 = c(x_1) - c(y_1)$$
- Similarly,

$$G_2 = c(x_1) - c(y_1) + c(x_2) - c(y_2).$$
- Gain criterion used by algorithm:
 Keep increasing k as long as $G_k > 0$.
- Note: this is a non-trivial addition because it allows for a temporary loss in gain:



- Neighbor limitation:



- LK limits the number of neighbors to the m nearest neighbors, where m is an algorithm parameter (e.g., $m=10$).
- Re-starts:
 - Recall: there are n choices for t_1 , the very first node.
 - LK tries all n before giving up.
- Best-tour: at all times LK records the best tour found so far.
- Note: LK is actually a little more complicated than described above, but these are the key ideas.

Performance:

- The standard heuristics (construction, K-OPT) give tours with 2-5% above Held-Karp.
- LK is usually between 1-2% off.


LKH-1: Lin-Kernighan-Helsgaun

From 1999-2009, Keld Helsgaun [Hels2009], added a number of sophisticated optimizations to the basic LK algorithm:

- The first set were added in 1999: [Hels1999].
 - We'll call this LKH-1.
- And the second set in 2009: [Hels2009].
 - We'll call this LKH-2.

Key ideas in LKH-1:

- Use $K=5$ (prefer this value of K over smaller ones).
 - Experimental evidence showed that the improvement going from 4- to 5-OPT is much better than 3- to 4-OPT.
 - Tradeoff: if K is too high, it takes too long
 - Fewer iterations
 - Less exploration of search space (even if you search a particular neighborhood more thoroughly).
- Relax *sequentiality* allow some x_i 's and y_i 's to repeat.
- Replace closest m neighbors with a different set of M neighbors:
 - Problem with LK:




- Recall best 1-tree in Held-Karp bound?

- Many of these edges are "good" edges for the tour.
- Experimental evidence: 70-80% of these edges are in optimal tour.
- LKH-1 idea: prefer 1-tree edges that go to neighbors.
- Let $L(T)$ = cost of best 1-tree
 - Can be computed fast (MST)
- For any edge e , let
 - $L(T, e)$ = cost of best 1-tree that *must* use e .
- How to force using an edge e ?
 - Find min-1-tree.
 - Add e to tree.
 - This causes a cycle.
 - Remove heaviest edge in cycle.
 - This leaves a min-1-tree that uses e .
- Define $\text{Alpha}(e) = L(T, e) - L(T)$ = importance of e in "1-tree-ness"
- Note: $\text{Alpha}(e) = 0$ for any edge in min-1-tree.
- LKH-1 sorts neighbors by α and uses best m of these.

LKH-2: Lin-Kernighan-Helsgaun, Part 2

Key additions to LKH-1:

- Allow K to increase beyond 5.
- Problem-partitioning:



- Divide points into clusters.
- Find best tour for each cluster.
- Stitch together into final tour.
- Run algorithm many times and merge "best parts" from multiple tours.
 - Called *iterative partial transcription*.
- Use sophisticated tour data structures to speed up running time.
- Results: million city problem with 0.058% of Held-Karp.
 - Within 0.058% of optimal.


Let's examine the partitioning idea:

- LKH-2 tries a number of partitionings, using different clustering algorithms.
- K-means clustering:

1. **repeat**
 - // Note: this is a different K than in K-OPT.
2. Pick k centroids.

- 3. Assign each point to closest centroid.
- 4. Re-compute the centroid based on assignments.
- 5. ~~until~~ no change

- Tour segmentation:
 - Run LKH-2 once to find a tour.
 - Segment the tour and re-solve the segments (partition).
- Geometric:









Recursive subdivision of space
(similar to k-d trees or quad-trees)

- Space-filling curve:

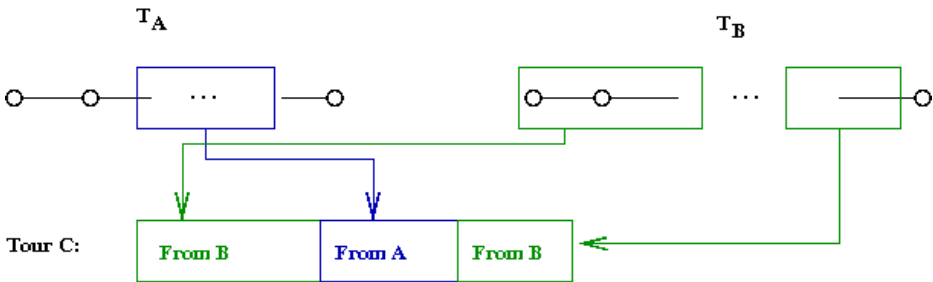
Hilbert space-filling curves (recursively defined)

Points close by along curve
are likely to be clustered




Iterative partial transcription (IPT):

- This is an idea from [Mobi1999].
- Goal: given two tours T_A and T_B , compute T_C that is better than both T_A and T_B .



- A single IPT *trial-swap* between tours T_A and T_B to creates tours $T_{A'}$ and $T_{B'}$




- An IPT-iteration:
 - Identifies all possible valid swap segments.
 - Tries the swaps and identifies the best possible tour that can be generated.
- How to use IPT:
 - Generate m tours T_1, \dots, T_m .
 - For each pair of tours i, j , perform an IPT-iteration.

Data structures

Given a K-OPT move, is the resulting "tour" a valid tour?


- Example:



- Naive way: walk along new tour T' to see if all vertices are visited
 → $O(n)$ per trial edge-swap
- Another problem: how to maintain tours?

Operations on tour data structures:


- First, note that any single swap can result in reversing the tour order for one of the segments affected:



- A single 2-OPT move will be called a *flip* operation.
- Also, any K-OPT move can be implemented by a sequence of 2-OPT moves.
→ LK-MOVE can be written to use *flip* operations.
- Other operations that need to be supported:
 - $next(a)$: the next node in tour order.
 - $prev(a)$: the previous node in tour order.
 - $between(a,b,c)$: determine whether b is between a and c in tour-order.
- Note: If a flip is performed correctly, it will result in a valid tour.
- Fredman et al. [Fred1995] show a lower bound of $(\log n) / (\log \log n)$ for these operations.

Arrays:


- Simple to implement.
- But consider what needs to be done to reverse a segment:



→ Can take $O(n)$.

Doubly-linked lists:


- flip* takes $O(1)$ pointer manipulations.
- Order reversal is also easy (comes for free): $O(1)$.




- But finding elements is hard: $O(n)$.

Modified splay trees:

- What is a splay tree?
 - Also called a *self-adjusting binary tree*.
 - [See lecture in algorithms course.](#)
 - Recall problem with binary trees: can go out of balance.
 - Problem with forced balance (e.g. AVL): too much overhead.
 - But use of *rotations* is useful.
 - Example of a splay-step: two mini-rotations:




- Another example:




- In a splay-tree: every accessed node is *splayed to the root*.
 - Similar to Move-to-Front in linked lists.

- Using a splay-tree for a tour:
 - Each node represents a city.
 - Initially, for first tour: in-order traversal is the tour:




Exercise: What is the tour represented by the above tree?

- Reversals are noted by marking intermediate nodes, e.g.




Tour: 8 5 4 3 6 7 1 2

- Each time a reversed-node is encountered, switch order (left swapped with right) in in-order traversal:




- Maintain an external array of pointers into tree, one per node.
- Implementing *next(a)*:
 - Recall *next(a)* in ordinary binary trees: leftmost node of the right subtree.



- Locate *a* using pointer-array: $O(1)$.

- Splay to root.
- Find successor using *tour-order* (instead of numeric order).
 - With no reversals, this is the leftmost node of the right subtree.
- With reversals, need to change direction for each flip (when recursing).
- The most complex operation is *flip()*:
 - Just like the splay-tree, there are several different cases.
 - Many involve some type of reversal.
 - The general idea (an example):



The segment tree:

- Devised by Applegate and Cook.
- Based on key observation about LK:
 - You try a sequence of flips (the LK-move).
 - When it doesn't work, you discard the whole sequence.
- In the data structures so far:
 - Every flip changes the data structure.
 - To discard, we need to *undo* flips in reverse order.
- A segment-tree tries to avoid the *undo* part.
 - Array representation of tour.
 - An auxiliary segment-list:
 - To help with tentative flips.
 - An auxiliary segment tree:
 - To help with fast navigation.

Performance:

- Segment-tree is usually best.
- 2-level list is next.
- Splay tree next (with theoretically the best performance).

Exact solution techniques: background

The general idea:

- Formulate TSP as a Integer Programming (IP) problem.
- Apply the *cutting-plane* approach.
- Judicious choice of cutting-plane heuristics.

But, first, what is Integer Programming? We'll need some background in *linear programming*.

Linear programming:

- The word *program* has different meaning than we are used to.
→ More like a "programme" of events.
- An LP (Linear Programming) problem is (in standard form):

$$\begin{array}{ll}
 \max & c_1x_1 + c_2x_2 + \dots + c_nx_n \\
 \text{such that} & \\
 & a_{11}x_1 + \dots + a_{1n}x_n \leq b_1 \\
 & a_{21}x_1 + \dots + a_{2n}x_n \leq b_2 \\
 & \cdot \\
 & \cdot \\
 & \cdot \\
 & a_{n1}x_1 + \dots + a_{nn}x_n \leq b_n \\
 \text{and} & \\
 & x_i \geq 0, \quad i=1, \dots, n \\
 & x_i \in \mathbb{R}
 \end{array}$$

- In vector/matrix notation:

$$\begin{array}{ll}
 \max & c^T x \\
 \text{s.t.} & Ax \leq b \\
 & x \geq 0
 \end{array}$$

- Example:

3.6 = selling price per unit of product 1

Amount of product 2 = x_2


$$\begin{array}{ll}
 \max & 3.6x_1 + 14x_2 \\
 \text{s.t.} & 12x_1 + 23.47x_2 \leq 400 \\
 & 34x_1 + 16x_2 \leq 650 \\
 & 55x_1 \leq 1200 \\
 & \text{and } x_i \geq 0
 \end{array}$$

only product 1 needs raw material #3

650 = total amount of raw material #2


16 = amount of raw material #2 used per unit of product 2

- Geometric intuition of inequality constraints ($Ax \leq b$):




- Each inequality defines a half-plane (half-space).
- The intersection is a polytope (polygon in 2D).
- The feasible region is sometimes called the *simplex*.

- If we plot objective function "lines":



- If we make a line-equation out of the objective function, some lines will pass through the feasible region.
- Clearly, we want the line with the highest "value" (for a max problem).
- Sweeping the line upwards (higher value), we want the line that is the last line to intersect the feasible region.
- This line always intersects the region at a *corner*.
- Three key algorithms, all major milestones in the development of LP:
 - George Dantzig's Simplex algorithm (1947).
 - Leonid Khachiyan's ellipsoid method (1979).
 - Narendra Karmarkar's interior-point method (1984).

- The simplex method:



- Start at a corner in the feasible region.
- A *simplex-move* is a move to a neighboring corner.
- Pick a better neighbor to move to (or even best neighbor).
- Repeat until you've reached optimal solution.
- What's known about the simplex method:

- Guaranteed to find optimal solution.
- Worst-case running time: exponential.
- In practice, it's quite efficient, approximately $O(n^3)$.
- Very efficient implementations available, both commercial and open-source.
- Has been used to solve very large problems (thousands of variables).
- What's known about the other algorithms:
 - Khachiyan's ellipsoid method: provably polynomial, but inefficient in practice.
 - Karmarkar's algorithm: provably polynomial and practically efficient for many types of LP problems.
- Note: an LP problem with equality constraints

$$\begin{array}{ll}\max & c^T x \\ \text{s.t.} & Ax = b \\ & x \geq 0\end{array}$$

can be converted to an equivalent one in standard form (with inequality constraints).

- Similarly, a min-problem can be converted to a max-problem.

Integer programming:


- An integer program (IP) is an LP problem with one additional constraint: all x_i 's are required to be integer:

$$\begin{array}{ll}\max & c^T x \\ \text{s.t.} & Ax \leq b \\ & x \geq 0 \\ & x \in \mathbb{Z}\end{array}$$

Exact solution techniques: TSP as an IP problem

First, let's express TSP as an IP problem:


- We'll assume the TSP is a Euclidean TSP (the formulation for a graph-TSP is similar).
- Let the variable x_{ij} represent the directed edge (i,j) .
- Let $c_{ij} = c_{ji}$ = the cost of the undirected edge (i,j) .



- Consider the following IP problem:

$$\begin{array}{ll}\min & \sum_{i,j} c_{i,j} x_{i,j} \\ \text{s.t.} & \sum_j x_{i,j} = 1 \quad // \text{ Only one outgoing arc from } i \\ & \sum_i x_{i,j} = 1 \quad // \text{ Only one incoming arc at } j\end{array}$$


- Unfortunately, this is not sufficient:



You can get multiple cycles.

→ Called *sub-tours*

- What to do? Consider this idea:



- Consider a subset of vertices S .
- In a valid tour,

$$\sum_{i,j} x_{ij} \leq |S| - 1 \text{ for all } i,j \in S.$$
- This is an inequality constraint that could be added to the IP problem.

→ Called a *sub-tour* constraint.
- How many such constraints need to be added to the IP problem?

→ One for each possible subset S .


→ Exponential number of constraints!
- Fortunately, one can add these constraints only as and when needed (see below).

Solving the IP problem:


- Naive approach:
 - Solve the *LP relaxation* problem first.

→ Remove integer constraints (temporarily) to get a regular LP, and solve it.
 - Round LP solution to nearest integers.


Unfortunately, this may not yield a feasible solution:




- Branch-and-bound:
 - We'll explain this for 0-1-IP problems (variables are binary-valued).
 - First, consider a simple exhaustive search, organized as a tree-search (the "branch" part):



- The tree itself can be explored in a variety of ways:
 - Breadth-first (high memory requirements)
 - High memory requirements.
 - Depth-first
 - Low memory requirements.
 - Cost-first
 - Expand the node that adds the least overall cost to the (partial) objective function.
- Note: if the cost to a node already exceeds the best tour so far, there's no need to explore further.
 - Parts of the tree can be *pruned*.



- Cutting planes:



- Add constraints to force the LP-solutions towards integers.
- With a sequence of such constraints, such a process can converge to an integer solution.
- However, it can take a long time.
- Gomory's algorithm:
 - A general cutting-plane algorithm for any IP.
 - The idea:
 - Solve LP.
 - Examine equations satisfied at corner point (of LP).
 - Round to integers in inequalities involving those variables.
 - Add these to constraints.
 - Repeat.
 - Unfortunately, it is slow in practice.

History of applying IP to TSP:


- Original cutting plane idea due to Dantzig, Fulkerson and Johnson in 1954.
 - Idea:

```

repeat
  solve LP
  identify sub-tours (cycles) and add corresponding "ISI-1" constraints.
until full-tour found

```

- Dantzig et al added a few more "sub-tour" like constraints.
- Today, there are several families of cutting-plane constraints for the TSP.
- Branch-and-cut
 - Cutting planes "ruled" until 1972.
 - Saman Hong (JHU) in 1972 combined cutting-planes with branch-and-bound
 - Called branch-and-cut.
 - The idea: some variables might change too slowly with cutting planes
 - For these, try both 0 and 1 (branch-and-bound idea).
 - Alternate way of viewing this:



- More sophisticated "cut" families:
 - Grotschel & Padberg, 1970's.
 - Padberg and Hong, 1980: 318-city problem.
 - Grotschel and Holland, 1987: 666-city problem.
 - Padberg and Rinaldi, 1987-88: combined multiple types of cuts, branch-and-cut and various tricks to solve 2392-city problem.
- During this time, LP techniques improved greatly
 - Can cut down "active" variables in an LP problem.
- Applegate et al (2006)
 - Sophisticated LP techniques, new data structures.
 - 85,900 city problem.

References and further reading

- [WP-1] [Wikipedia entry on TSP.](#)
- [WP-1] [Georgia Tech website on TSP.](#)
- [Appl2006] D.L.Applegate, R.E.Bixby, V.Chvatal and W.J.Cook. *The Traveling Salesman Problem*, Princeton Univ. Press, 2006.
- [Aror1992] S.Arora, C.Lund, R.Motwani, M.Sudan and M.Szegedy. Proof verification and hardness of approximation problems. *Proc. Symp. Foundations of Computer Science*, 1992, pp.14-23.
- [Aror1998] S.Arora. [Polynomial Time Approximation Schemes for Euclidean Traveling Salesman and other Geometric Problems](#). *J.ACM*, 45:5, 1998, pp. 753-782.
- [Bear1959] J.Beardwood, J.H.Halton and J.M.Hammersley. The Shortest Path Through Many Points. *Proc. Cambridge Phil. Soc.*, 55, 1959, pp.299-327.
- [Chan1994] B.Chandra, H.Karloff and C.Tovey. [New results on the old k-opt algorithm for the TSP](#). *5th ACM-SIAM Symp. on Discrete Algorithms*, 1994, pp.150-159.
- [Clar1964] G.Clark and J.W.Wright. Scheduling of vehicles from a central depot to a number of delivery points. *Op.Res.*, 12, 1964, pp.568-581.
- [Chri1976] N.Christofides. Worst-case analysis of a new heuristic for the travelling salesman problem. Report No. 388,

GSIA, Carnegie-Mellon University, Pittsburgh, PA, 1976.

[Croe1958] G.A.Croes. A method for solving traveling salesman problems. *Op.Res.*, 6, 1958, pp.791-812.

[Fred1995] M.L.Fredman, D.S.Johnson, L.A.McGeogh and G.Ostheimer. Data structures for traveling salesmen. *J.Algorithms*, Vol.18, 1995, pp.432-479.

[Glov1990] F.Glover. Tabu Search: A Tutorial, *Interfaces*, 20:1, 1990, pp.74-94.

[Guti2007] G.Gutin and A.Yeo. The Greedy Algorithm for the Symmetric TSP. *Algorithmic Oper. Res.*, Vol.2, 2007, pp.33--36.

[Held1970] M.Held and R.M.Karp. The traveling-salesman problem and minimum spanning trees. *Op.Res.*, 18, 1970, pp.1138-1162.

[Hels1998] K. Helsgaun. An Effective Implementation of the Lin-Kernighan Traveling Salesman Heuristic, DATALOGISKE SKRIFTER (Writings on Computer Science), No. 81, 1998, Roskilde University.

[Hels2009] K. Helsgaun. General k-opt submoves for the Lin-Kernighan TSP heuristic. *Mathematical Programming Computation*, 2009.

[John1997] D.S.Johnson and L.A.McGeoch. The Traveling Salesman Problem: A Case Study in Local Optimization. In Aarts, E. H. L.; Lenstra, J. K., *Local Search in Combinatorial Optimisation*, John Wiley and Sons Ltd, pp. 215;V310, 1997.

[Karp1972] R.Karp. Reducibility among combinatorial problems, in R. E. Miller and J. W. Thatcher (editors). , New York: Plenum. pp. 85-103.

[Kirk1983] S.Kirkpatrick, C.D.Gelatt, and M.P.Vecchi. Optimization by Simulated Annealing. *Science*, 220 1983, pp.671-680.

[Lin1973] S.Lin and B.W.Kernighan. An Effective Heuristic Algorithm for the Traveling- Salesman Problem. *Op.Res.*, 21, 1973, pp.498-516.

[Mobi1999] A.Mobius, B.Freisleben, P.Merz and M.Schreiber. Combinatorial optimization by iterative partial transcription. *Phys.Rev. E*, 59:4, 1999, pp.4667-74.

[Rein1991] G.Reinelt. TSPLIB :X A Traveling Salesman Problem Library. *ORSA J. Comp.*, 3:4, 1991, pp. 376-384.

D.J.Rosenkrantz, R.E.Stearns and P.M.Lewis. An analysis of several heuristics for the traveling salesman problem. *SIAM J. Computing*, Vol.6, 1977, pp.563-581.

[Sahn1976] S.Sahni and T.Gonzalez. P-complete approximation problems. *J.ACM*, Vol.23, 1976, pp.555-565.

[CS153] R.Simha. Course notes for CS-153 (Undergraduate algorithms course).

[Vale1997] C.L.Valenzuela and A.J.Jones. Estimating the Held-Karp lower bound for the geometric TSP. *European J. Op. Res.*, 102:1, 1997, pp.157-175.

Note: The Hilbert curve was an image found on Wiki-commons.
